

Tamad-Strat

Quant validation report — the 3-candle reversal, Revelio-style


Data: Binance spot 2017-08-01 → 2025-01-01 · BTC / ETH / SOL · 5m / 15m / 1h · 8,784 final-config trades

Train $\leq$ 2022 · Test 2023–2024 · HOLDOUT 2025+ still locked (sign-off pending on issue #18)

VERDICT

1. The naked pattern is a coin flip — 97,535 trades, win rate pinned at the 25% breakeven.
2. Every taught rule is inert or harmful: the ideal sweep, zone confluence stacking, HTF bias.
3. A real pre-cost edge exists, found by ablation: $C1 \geq 1 \times \text{ATR} + \text{order-block zone} + \text{US session}$, exited with a 1R trailing stop — 9/10 training picks confirmed on unseen 2023–24 data.
4. Spot taker costs destroy it: tight pattern stops make 0.22% round trips cost 0.2–1.4R per trade.
Every with-costs portfolio is negative. The 60% claim is not in the written rules.

The claim vs the data




At 1:3 RR the breakeven win rate is 25%; 60% over 300 trades would mean +1.4R per trade ($\approx +420R$ total), far outside the noise band (95% CI $\approx \pm 5.5pp$). Nothing mechanical in the written rules reproduces it. The trailing exit's higher win rate is not comparable to the 25% line — its payoff distribution differs (many small +0.1R wins). The 60% almost certainly lives in discretion or journaling, not in the rules as written.

The version ladder — every rule made to earn its place


stage	what was tested	verdict	source (docs/)
V0 naked pattern	all triplets, 1:3, no filters	coin flip — WR = breakeven, PF $\approx$ 1.0	V0_RESULTS.md
V1 ideal sweep	C2 sweeps C1 low/high	no lift; bottom of ablation ranks	PILLS_RESULTS.md
V1 C1 $\geq 1 \times$ ATR	climax first candle	modest lift — first real clue	PILLS_RESULTS.md
V2 as taught	pattern at any of 5 zone types	filters nothing (~70% qualify)	ZONES_RESULTS.md
V2 confluence stack	2+ zones agree	flat: PF 0.98 / 1.01 / 1.00	ZONES_RESULTS.md
V3 HTF bias	4h/1d pattern gates LTF entries	hurts at every bias age	V3_RESULTS.md
Ablation 336 configs	train-select → test-confirm	survivor: C1 $\geq 1 \times$ ATR + OB + US (9/10 held)	ABLATION_RESULTS.md
Exits	1:3 vs trail 1R/0.5R vs ATR-trail	trail_1R best OOS (test PF 1.44 vs 1.18)	EXITS_RESULTS.md
Risk	fixed vs HWM vs parity sizing	parity sizing holds OOS	RISK_RESULTS.md
Costs	0.22% notional round trip	ALL portfolios negative; ETH 1h $\approx$ breakeven	COSTS_RESULTS.md

Final configuration on 5m — equity in R (flat 1R per trade)




Final configuration on 15m — equity in R (flat 1R per trade)


BTCUSDT — 689 trades · PF 1.43 zero-cost → 0.33 with costs




ETHUSDT — 680 trades · PF 1.49 zero-cost → 0.46 with costs




SOLUSDT — 373 trades · PF 1.50 zero-cost → 0.68 with costs



Final configuration on 1h — equity in R (flat 1R per trade)




The kill-shot: profit factor before vs after spot taker costs



Why: risk per trade is the tiny distance to the pattern-extreme stop, but cost is charged on the full notional — notional/risk ratios of 300-650× turn a 0.22% round trip into 0.7-1.4R on 5m. The pattern's defining feature (the tight stop) is exactly what makes it cost-doomed. A BNB fee discount (~23% off) changes no signs. Only ETH 1h survives at ≈ breakeven (PF 1.04).

Parity portfolios on the TEST window (calibrated on training only)

High-watermark sizing, per-asset risk bisected on training trades to a 20% standalone max drawdown, then applied unseen to 2023–2024. \$10,000 start.



Log scale. The blue curves are the pre-cost edge (real, confirmed OOS). The red curves are what a spot taker account would actually experience.

What was optimized — and what could still move the needle

Optimization already done (the Revelio ladder)

Pills one at a time (sweep, C1 size) · five zone families alone and stacked · HTF bias at every age · 336-config ablation grid selected on training and confirmed on unseen 2023–24 · four exit structures · three sizing schemes with drawdown-parity calibration · costs applied last. Result: $PF \approx 1.0 \rightarrow 1.44$ (test, pre-cost). The strategy that survived shares only the 3-candle trigger with the one taught.

The binding constraint is execution cost, not entry quality

Further entry filters cannot fix a strategy paying 0.2–1.4R per trade in fees. The levers that remain:

1. Wider stops

Stop at pattern-extreme minus an ATR buffer (or the zone edge) shrinks the notional/risk ratio and directly divides the cost in R. Test whether the pre-cost edge survives the wider stop.

2. Maker / futures execution

Binance futures maker is 0.02%/side (~10× cheaper). Entry as a limit at C3 close with conservative fill modeling (require trade-through). Different strategy — fills are uncertain — but the only path that plausibly flips the after-cost sign on low timeframes.

3. Trade where costs are survivable

1h already pays only ~0.2R/trade; 4h as a traded timeframe (not just bias) is untested. Widen the asset basket at 1h+ to rebuild trade count.

4. What the ML layer (#19) already settled

A random forest over 28 pre-trade features (Revelio Phase 7) confirmed OOS: the edge concentrates in the TIGHTEST structures (risk% filter, +12–23% test PF on every interval) — so the cost-doom is intrinsic, and widening stops trades away edge. One real refinement: a clean hold candle (small front wick), the only rule that also improved after-cost PF (1h: 0.73 → 0.88, still losing). docs/ML_RESULTS.md.